

USAII GLOBAL AI HACKATHON 2026
COMMUNITY TRACK

CLEARPATH AI

WE DON'T JUST FIND HELP.
WE SHOW YOU WHY.

THE PROBLEM

When someone needs help, a confident wrong answer is more dangerous than no answer at all.

Someone types 'I can't pay rent.' A search engine gives rental assistance programs. An AI confidently recommends a 3-month waitlist. But what they actually needed was legal aid for an emergency eviction defense — help available within 24 hours.

The confident match just cost them their home.

73%

OF COMMUNITY RESOURCE SEARCHES LEAD TO DEAD ENDS

THE INSIGHT

Most AI systems hide their uncertainty. In community resource matching, that is dangerous.

What if the AI showed you what it knows, what it doesn't, and why?

When someone is looking for help, an honest uncertain answer can save a life. A confident wrong answer can destroy one. Uncertainty is not a bug — it is the safety feature.

THE SOLUTION

CLEARPATH AI

A community resource navigator that shows calibrated confidence instead of hiding uncertainty.

Six layers of intelligent transparency:

1. Free-text input — Describe your situation in your own words
2. Crisis safety net — Hardcoded keyword detection, AI cannot override
3. Multi-label classification — Multiple matches with confidence scores, not one answer
4. Clarification questions — When confidence is low, ask — don't guess harder
5. Transparent reasoning — Every recommendation shows why, what else, and how confident
6. Human escalation — When AI is unsure, it backs off and connects to a real person

HOW IT WORKS

1

User de- scrib- es need

Free-
text: 'I
can't
pay
rent
any-
more'

2

Crisi- s chec- k

Keyw-
ord
scan
trig-
gers
211
hot-
line
dis-
play

3

NLP clas- sifi- cation

Zero-
shot
multi-
label:
Legal
Aid
78%,
Emer-
gency
Funds
52%,
Housi-
ng
34%

4

Clarif- ica- tion

AI PIPELINE

- Hugging Face Transformers — Zero-shot NLP classification
- Temperature scaling — Confidence calibration
- Active learning — Clarification question logic

WEB APPLICATION

- React + TypeScript — Frontend interface
- Node.js + Express — Backend API
- MongoDB — Resource database

'Have you received an eviction notice?'
— Yes
→ Legal Aid
94%

5

Transparent output

3 resources per category with reasoning and confidence

6

Human escalation

Low confidence
→ connect to navigator at 211

RESPONSIBLE AI

- Python key-word engine

— Crisis de-

RESPONSIBILITY IS THE ARCHITECTURE, NOT A DISCLAIMER.

tection layer

Three built-in safety mechanisms:

- Vercel + AWS — Free-tier hosting
- 1. Crisis safety net — Hardcoded keyword detection that the AI layer CANNOT override. Crisis resources appear regardless of classification results.
- 2. Forced transparency — Confidence scores are always visible.
- No paid APIs required
- 3. Automatic human escalation — When no category exceeds 70% confidence, the system recommends speaking with a trained navigator instead of guessing harder.

KNOWN LIMITATIONS

- Cannot determine eligibility for programs
- Cannot replace professional social workers
- May misclassify complex multi-need situations
- Requires curated, locally accurate resource data

**WHEN IT
MATTERS MOST,
HONESTY IS
THE SAFEST
ANSWER.**

CLEARPATH AI
BUILT BY AMINE & HARSHIT

